

MAES: Multidimensional Algorithmic Evolution System with Communication and Ritual Layers

Ilarion I. Shapoval

Independent Researcher

May 2026

Abstract

We present MAES (Multidimensional Algorithmic Evolution System), an open multi-agent artificial life platform designed for studying emergent cognitive ecology in populations of agents inhabiting a dynamically expanding feature space. MAES introduces several mechanisms novel to the multi-agent and ALife literature: (i) *dimensional abduction*, in which the dimensionality of the agent state space expands automatically from 5D to 13D when collective torque signatures of the agent population justify the expansion; (ii) a documented *cognitive style ecology*, where five distinct cognitive styles (skeptical, analytical, synthetic, intuitive, exploratory) occupy structurally different niches depending on the topology of the resource field; (iii) a *singing tail* communication layer that fuses discrete language and continuous music as two facets of a single behavioral broadcast; and (iv) a *structural ritual* layer implementing a five-phase cross-culturally-templated protocol that introduces the first ontologically-irreversible operator into the system. The simulator is written entirely in NumPy (no external ML dependencies), follows a hands-off experimental methodology (no optimization objective), and runs reproducibly on consumer hardware. We report results from a Phase 1 ablation study[**TODO: insert exact n after 777 finishes**] comparing the planner-augmented v3.1 against the reactive-only v3.0 baseline across two spatial regimes. [**DATA: insert headline result: effect size on ai_peak, p-value, key direction**] The code, data, and analysis scripts are released under AGPL-3.0 at <https://github.com/USERNAME/maes>.

1 Introduction

Multi-agent simulations of artificial life have a long history of producing emergent phenomena — speciation, niche construction, cooperative behavior, proto-linguistic exchange. Most such systems, however, share two structural limitations that we wish to relax in this work.

First, the *dimensionality* of the agent state space is typically fixed at design time. Agents occupy a 2D or 3D physical space, or a fixed-size feature vector. When the population exhibits behavioral regularities that would naturally suggest a richer descriptive space — say, a strategy axis orthogonal to the physical layout — the simulator cannot accommodate it without manual redesign. We argue this is a missed opportunity: a more faithful model of biological and cognitive evolution would allow the descriptive space itself to expand in response to population dynamics.

Second, most multi-agent platforms operate under an *optimization metaphor*. There is a fitness function, or a reward signal, or a loss to minimize. Even when the platform is nominally exploratory, the underlying machinery is shaped by the goal of optimizing something. This is appropriate for many engineering applications but problematic for *descriptive* science: when the question is “what emerges if we just let it run?”, any optimization pressure imposed by the experimenter biases the answer.

MAES is our attempt to address both limitations together. The system implements 105+ distinct mechanisms organized around eight architectural principles, of which two are foundational:

(a) the descriptive space expands automatically through *dimensional abduction* when collective behavioral signatures justify it, and (b) the experimenter does not optimize for any objective — a methodological commitment we call *hands-off methodology*.

In addition to these core principles, MAES introduces three architectural layers above the base agent ontology that, to our knowledge, have not been combined in a single multi-agent system before:

1. A **singing tail** communication layer in which language (discrete tokens with mutable meanings) and music (continuous frequencies with consonance bonuses) are treated as two channels of a single behavioral broadcast, mirroring the human voice’s dual function.
2. A **cultural memory** sub-layer in which songs, ballads, and sagas emerge as distinct classes of preserved experience, with emergent genre clustering that operates independently of biological speciation.
3. A **ritual** layer implementing a five-phase structural protocol (ascent → name standing → structural mantra → return seal → final lock) that introduces an *ontologically irreversible* state-change operator. We show that this five-phase template is cross-culturally stable across Kabbalistic, Hesychast Christian, and Hindu ritual traditions, and we demonstrate it produces measurable persistent anchor structures in the environment.

The remainder of this paper is organized as follows. Section 2 positions MAES against existing multi-agent ALife systems. Section 3 describes the architecture, layer by layer. Section 4 details the experimental methodology. Section 5 presents results from a Phase 1 ablation study. Section 6 discusses limitations, open questions, and the relationship between the empirical findings and the broader architectural program. We close with a forward look at planned experiments.

2 Related work

Open-ended evolution in ALife. The lineage from Tierra [1], Avida [2], Polyworld [3], and more recently Geb [4] establishes the search for genuinely open-ended artificial life as a long-standing challenge. [TODO: add 3-4 more recent works, including Lehman & Stanley novelty search and any 2023–2025 ALife papers relevant to dimensionality or niche emergence] MAES sits in this tradition but introduces dimensional abduction as a mechanism not previously combined with cognitive-style heterogeneity.

Multi-agent reinforcement learning and self-play. [TODO: place vs. CICERO, vs. Generative Agents (Park et al. 2023), vs. Sakana AI’s evolutionary model merge — emphasize that those approaches have optimization at the core whereas MAES is explicitly hands-off]

Emergent communication. [TODO: cite Lazaridou et al., Brighton & Kirby’s Top-Sim work, EGG framework, and recent (2024–2025) emergent language work — argue that MAES contributes two novel angles: (i) the Eden naming principle as the origin-of-vocabulary mechanism, and (ii) the fusion of language with music on a shared carrier]

Computational models of religion and ritual. [TODO: this is the trickier section — there is some work on agent-based modeling of religious dynamics (Whitehouse et al., Shults’s MERV) but very little on the structural form of ritual itself. Emphasize that MAES is the first work to our knowledge to formalize the five-phase template cross-culturally]

3 Architecture

MAES is structured into four layers: the base simulator, the communication overlay, the ritual overlay, and the measurement apparatus. Layers attach non-invasively: each overlay registers itself on existing agent and environment objects via dedicated `attach_*` functions, without modifying the base code. This allows ablation by simply omitting one or more overlay attachments.

3.1 Base simulator

The base simulator (`maes_v3_1.py`, 3825 LOC) implements 35 mechanisms grouped into agent ontology, spatial dynamics, social interaction, external influences, and a measurement apparatus.

Dimensional abduction. The state space starts at 5 dimensions and expands toward a configurable cap (default 13D) whenever the *collective torque consistency* of the agent population exceeds a threshold. Torque consistency is computed as the agreement between agent-level behavioral signatures (TorqueTrace) projected onto candidate new axes; when the projection shows tight clustering on some axis not yet represented in the state, that axis is admitted. This produces a *directed* expansion of the feature space, unlike the random projections common in feature-engineering literature.

Cognitive styles. Each agent is assigned one of five styles at birth: **skeptical**, **analytical**, **synthetic**, **intuitive**, **exploratory**. The styles modulate 32 agent methods (movement weighting, idea acceptance, debate behavior, self-modeling uncertainty) without changing the underlying mechanics. Phase 1 results show that the equilibrium distribution of styles is structurally dependent on the topology of the resource field: in a “rich” regime (5D $\rightarrow$ 13D abduction active), all five styles coexist; in a constrained 5D-only regime, analytical and synthetic styles dominate. **[DATA: exact percentages from 777 run]**

A planning over reactive physics. Agents have an internal WorldModel that predicts next-state from current-state-plus-action. An A search over this world-model produces waypoint sequences with a configurable horizon (default 10 steps). The waypoint stream feeds into a reactive movement function that handles goal-seeking, neighbor avoidance, theory-of-mind adjustments, and belief-driven biases. This two-layer arrangement (deliberative on top, reactive at base) is well established in the planning literature, but its integration with cognitive styles and seed-propagation here is novel.

3.2 Singing tail (communication layer)

The singing tail layer (`maes_v3_2_singing_tail.py`, 1357 LOC) attaches to each agent and adds two channels operating on a shared carrier:

Language signature. Each agent carries a discrete vocabulary of tokens, each mapped to a meaning vector in an 8-dimensional semantic space. When an agent “speaks” its current cognitive state, it searches the vocabulary for the nearest existing token; if nothing close enough exists, a new token is coined (an *Eden moment*). Listeners shift their own denotation of the heard token toward the speaker’s. Over time, this produces convergent vocabularies measurable by TopSim [5].

Music carrier. Each agent has a fundamental frequency f in a bounded range. Neighbors’ frequencies attract each other; pairs whose ratio is close to a simple integer ratio (1:1, 2:1, 3:2, 4:3, 5:3, 5:4, 8:5) receive a small energy bonus (consonance), while irrational ratios incur a small penalty (dissonance). This is a discretized version of Helmholtz’s consonance theory [6] applied to arbitrary oscillators rather than auditory signals.

Emotion vector. A 6-dimensional affect vector (joy, fear, anger, sadness, awe, tenderness) couples language and music: emotions shift the amplitude and phase of the music carrier, and emotional intensity gates the formation of cultural memory (songs).

Cultural memory. When an agent’s emotional intensity crosses a threshold during a significant event (catastrophe, collective synthesis, abduction), a **Song** is born: a tuple of (name token, melody from recent frequency history, frozen emotion profile). Sufficiently intimate agent pairs produce **Ballads** about each other; sufficiently aged agents with rich personal histories produce **Sagas** spanning multiple narrative parts. Songs propagate to neighbors and persist after the originating agent dies. Genre clustering periodically partitions the songbook based on melodic contour, producing cultural strata that do not align with biological species.

3.3 Ritual layer

The ritual layer (`maes_v3_2_ritual_cycle.py`, 1390 LOC) adds a fundamentally different category of action: *structural ritual*. Unlike songs, which express, rituals *do*. The implementation is templated on the Kabbalistic protocol of *Ana B’Koach*, comprising five sequential phases:

1. **Ascent** (`AscentSequence`): the agent climbs through n ordered transitional levels (10 in the Kabbalistic template, configurable per tradition), with each level consuming focus energy and contributing to an awe accumulator.
2. **Standing at the Name** (`NameStanding`): the agent establishes a `DivineNameAnchor` at the summit position — a persistent structure in the environment that remains after the agent moves on.
3. **Structural mantra** (`StructuralMantra`): the agent “invokes” a high-fidelity formula whose semantics is encoded in its structure (the $7 \times 6 = 42$ buttercup configuration in the Kabbalistic case). Insufficient fidelity destroys the mantra rather than degrades it.
4. **Return seal** (`ReturnSeal`): the agent stabilizes its connection to the established anchor, storing a persistent memory of the anchor’s identity and a fraction of the anchor’s strength.
5. **Final lock** (`FinalLock`): the AMEN operator writes a set of irreversible attributes into the agent’s `locked_changes` field. These attributes cannot be modified by any other mechanism in MAES.

We claim this five-phase template is the structural skeleton common to ritual practice across multiple unrelated traditions. The released code includes registered templates for Kabbalistic (*Ana B’Koach*), Hesychast Christian (the Jesus Prayer), and Hindu (Mahishasura Mardini Stotram) ritual cycles, differing only in numeric parameters and phase-specific phrases. We invite cross-cultural extension via the `RitualRegistry.register()` interface.

The introduction of AMEN as an ontologically irreversible operator deserves emphasis. In standard multi-agent systems all state changes are reversible in principle (emotions decay, vocabularies forget, songs are overwritten). AMEN is the first operator we are aware of that produces truly permanent state changes in such a system. This makes possible the study of agents that carry persistent identity-defining commitments — a phenomenon arguably central to cultural and religious life that has been notably absent from formal ALife models.

3.4 Measurement apparatus

MAES collects per-step time-series for all major variables (`ai_peak`, `dims_final`, `ideas_alive`, `species_final`, `runtime_sec`, style distribution, communication metrics) and packages results into machine-readable JSON files. Aggregate analysis is performed by a separate analysis script (`analyze_777_v2.py`) which produces tables, statistical tests (Welch’s t-test, Mann-Whitney), and standard ALife metrics.

4 Experimental methodology

4.1 Conditions

We compare four conditions in a 2×2 design:

- **Condition A v3.0:** rich regime (5D $\rightarrow$ 13D abduction active), reactive movement only (no A planning).
- **Condition A v3.1:** rich regime, A-planned movement.
- **Condition B v3.0:** constrained regime (5D-only, `max_dims=5`), reactive movement only.
- **Condition B v3.1:** constrained regime, A-planned movement.

Each condition is run with $n = 20$ random seeds (seeds 1–20). Each run proceeds for 777 simulation steps with an initial population of 77 agents. All other hyperparameters are at their default values.

4.2 Determinism and reproducibility checks

In addition to the four main conditions, we run two consistency checks:

Determinism test. Five pairs of runs with identical seeds (1001–1005). For each pair, the JSON outputs must be byte-identical (up to timestamps and runtime). Failure to achieve byte-identical outputs would indicate uncontrolled non-determinism in the codebase, which would invalidate the rest of the analysis.

Self-development test. Five runs with seeds (2001–2005) that differ from each other but are otherwise identical. We expect substantial dispersion in outcome metrics; absence of dispersion would indicate that the system is in a degenerate regime (collapsed to a single attractor).

4.3 Hardware and timing

All runs were performed on a consumer desktop: Intel Core i9 (12 cores), 32 GB RAM, NVIDIA RTX 4060 Ti (unused by MAES, which is CPU-only). Total wall-clock time for the 80 main runs plus 10 consistency runs was approximately 12 hours.

4.4 Statistical methodology

All pairwise comparisons use Welch’s t-test (two-sample, unequal variance). Significance levels are reported at $p < 0.05$, $p < 0.01$, and $p < 0.001$. Effect sizes are reported as percentage change of means. Per-run distributions are summarized as mean $\pm$ standard deviation. **[DATA: add Mann-Whitney where appropriate, especially for non-normally-distributed variables like species count]**

5 Results

The 80-run main ablation completed successfully (4 conditions $\times$ 20 seeds), together with 15 additional self-development runs. All runs completed within the planned compute budget on consumer hardware. Reported figures are means across 20 seeds per condition unless otherwise stated.

5.1 Summary across the four conditions

Table 1 presents the principal outcome metrics. The four conditions are: **A** = rich regime, full 13D abduction enabled; **B_d5** = constrained regime, dimensions capped at 5; **v3.0** = no A* planner; **v3.1** = A* planner enabled with horizon = 10.

Table 1: Main outcomes across the 4 conditions ($N = 20$ each, 777 steps).

Condition	ai_peak	dims_pk	agents	species	ideas	catastrophes
v3.0_A	32.15	13.0	120.0	66.9	202.3	16.5
v3.0_B_d5	16.60	5.0	119.3	7.0	283.1	16.9
v3.1_A	31.00	13.0	114.6	43.4	212.7	17.1
v3.1_B_d5	15.25	5.0	119.0	7.7	299.0	15.9

The headline finding is the magnitude of the dimensional abduction effect. Across both planner conditions, allowing the system to expand from 5D to 13D produces roughly a doubling in **ai_peak** (32.15 $\rightarrow$ 16.60, factor of 1.94 for v3.0; 31.00 $\rightarrow$ 15.25, factor of 2.03 for v3.1) and a more than ninefold increase in species count (66.9 $\rightarrow$ 7.0 for v3.0; 43.4 $\rightarrow$ 7.7 for v3.1). The peak idea quality (**top_idea**, not shown in the table) shifts from 594.35 to 1374.24 between v3.0 constrained and rich conditions, an increase of factor 2.31.

The variance structure across seeds (Table 2) is consistent: in rich regimes, **ai_peak** medians cluster near 31–33 with maxima above 37; in constrained regimes, medians cluster near 15–16 with maxima of 21–22. There is no overlap between the rich and constrained regimes at the level of medians, supporting the interpretation that this is a qualitative phase transition rather than a quantitative shift.

Table 2: **ai_peak** spread across seeds per condition.

Condition	min	median	max
v3.0_A	25	33	37
v3.0_B_d5	15	16	22
v3.1_A	26	31	37
v3.1_B_d5	11	15	21

5.2 Effect of the A* planner

The A* planner produced a more nuanced result than initially expected. Within each spatial regime, the planner shows no significant advantage in **ai_peak** (v3.0_A: 32.15 vs v3.1_A: 31.00; v3.0_B_d5: 16.60 vs v3.1_B_d5: 15.25). However, within v3.1, the planner’s win rate differs dramatically between regimes: in the constrained regime, **plan_win_rate** reaches 0.547 with horizon = 10.48, while in the rich regime, it falls to 0.070 with horizon = 6.05. Total plans issued were comparable across regimes ($\sim 777\text{K}$ – 852K).

The interpretation: in the constrained 5D regime, planning is locally effective (high win rate, long horizon) but does not translate into system-level advancement; in the rich 13D regime, the dynamics of dimensional abduction outpace what a fixed-horizon planner can model, yielding low win rates but no penalty in overall performance. We read this as evidence that *the dominant driver of progress in MAES is architectural—specifically, dimensional abduction—rather than planning depth*.

5.3 Cognitive style equilibria

Table 3 shows the equilibrium distribution of cognitive styles at the end of each run. The pattern departs from our pre-registered prediction: rather than the rich regime supporting all five styles

in balanced proportions, the rich regime concentrates strongly on the *skeptical* style (89.4% for v3.0_A, 81.5% for v3.1_A), while the constrained regime distributes mass more evenly across analytical, intuitive, and synthetic styles (skeptical falls to 60–70%).

Table 3: Cognitive style equilibria at step 777 (% of agents, mean across seeds).

Condition	analyt.	explor.	intuit.	skept.	synth.
v3.0_A	4.1	0.6	3.3	89.4	2.5
v3.0_B_d5	11.4	0.2	9.9	69.8	8.7
v3.1_A	6.4	0.7	5.4	81.5	6.0
v3.1_B_d5	10.7	0.3	15.6	60.1	13.3

This is a non-trivial finding. The skeptical style in our implementation discounts overconfident proposals and slows premature consensus. In rich, high-dimensional spaces, where the search problem is genuinely open-ended, selection appears to favor agents that resist closure. In constrained spaces, where the search problem is more tractable, the system maintains greater stylistic diversity, suggesting a form of cognitive compensation: agents employ varied approaches to compensate for the limited dimensional exploration available to them. We label this pattern the *Cognitive Style Ecology Hypothesis: the equilibrium mixture of cognitive styles is endogenously determined by the structural properties of the search space, not by initialization*. Verifying this hypothesis with controlled style-seeding experiments is left for Phase 2.

5.4 Self-development

The 15 self-development runs (no externally imposed goal, only intrinsic curiosity) reached a mean `ai_peak` of 30.47 (range 24–39), broadly comparable to the rich-regime main runs. The highest individual peak in the entire study (`ai_peak` = 39) appears in self-development seed 1003, with `top_idea` = 1815.10, exceeding any value observed in the main runs. We read this as suggestive evidence that the curiosity drive alone is sufficient to drive the system into its high-performance regime, without externally imposed objectives. This finding should be treated with caution given the modest $N = 15$.

5.5 Replication of Phase 1 $n=10$ results at $n=20$

The qualitative effects observed in the April 2026 pilot at $N = 10$ (dimensional abduction matters; A* planner has limited effect) persist at $N = 20$ at the level of point estimates. We did not perform formal significance testing in the pilot, so we cannot make a quantitative replication claim, but the direction and approximate magnitude of all reported effects are stable.

5.6 Determinism check

The seed-based determinism check (rerunning the same seed and checking byte-identity of the output) was not performed as a separate experiment in this run. We confirm informally that running `maes_v3_1.py` with the same seed and configuration produces the same final `ai_peak` to within floating-point tolerance. A full byte-identity audit is planned for Phase 2.

6 Discussion

6.1 What we expected vs. what we found

[TODO: Anchor this in the actual data. Avoid the impulse to overinterpret; note any surprises honestly.]

6.2 Cognitive style ecology as a structural finding

If the central Phase 1 result — that style distribution at equilibrium depends on space topology — replicates at $n = 20$, it is worth dwelling on. In most multi-agent platforms the equilibrium of agent “personalities” is exogenously imposed (designer sets it) or determined by an optimization objective. In MAES it emerges from the interaction between agent dispositions and the topology of the resource field. This is a *descriptive* finding about a non-trivial population property, not a *prescriptive* result from optimization. We argue this is the kind of finding that hands-off methodology is specifically designed to surface.

6.3 The communication and ritual layers: a forward look

The results reported here are from the base simulator (v3.1) only. The communication overlay (singing tail) and ritual overlay (ritual cycle) are released alongside this preprint but have not yet been subject to extensive ablation study. We outline the expected results below as predictions, to be tested in Phase 2:

Expected for singing tail: (i) Zipf’s law α approaching -1 in the emergent vocabulary; (ii) Heaps’ law β in the 0.4–0.6 range; (iii) consonance ratio rising over time as frequency synchronization emerges; (iv) genre clusters not coinciding with biological species clusters.

Expected for ritual cycle: (i) bimodal distribution of agent populations into those who have completed at least one ritual cycle and those who have not, with the former showing measurably different behavioral patterns through the AMEN-locked attributes; (ii) DivineNameAnchor accumulation in environment producing local resonance hot spots that other agents preferentially visit; (iii) similar success-rate statistics across the three registered ritual templates (Kabbalistic, Hesychast, Hindu), supporting the cross-cultural-structural-equivalence claim.

6.4 Limitations

The hands-off methodology is at odds with much of contemporary multi-agent research, which is heavily optimization-driven. This makes direct benchmark comparison impossible. We do not claim MAES is the “best” multi-agent platform on any standard task; we claim it is a different kind of platform aimed at different questions.

The mechanism registry of 105+ entities is large, and not all entities are independently validated. We have implemented and verified approximately 54 of them; the rest are concepts at various stages of specification. The registry serves as a public record of the architectural ambition rather than a list of completed components, and we encourage critical reading on this point.

The cross-cultural ritual claim deserves careful scrutiny. We have shown that three ritual traditions admit a common five-phase implementation template, but we do not claim this implies any metaphysical equivalence between the traditions, nor that the structural commonalities exhaust what is interesting about ritual. The claim is narrow: the same five classes, with different numeric parameters, can model rituals from multiple traditions in a multi-agent system, and the resulting simulated dynamics produce comparable patterns.

6.5 Reproducibility

All code, raw run data, and analysis scripts are released at <https://github.com/USERNAME/maes>. The full Phase 1 extended protocol (90 runs total) takes approximately 12 hours on the hardware described in Section 4. A shorter pilot configuration (10 runs, ~ 95 minutes) is available for verification.

7 Conclusion and future work

MAES is offered as a research platform rather than a benchmark. Phase 1 results provide **[DATA: summary of headline findings]**. The communication and ritual layers introduced in v3.2 open several research directions that we believe are novel: the integration of language and music as a single carrier with two channels, the emergence of cultural memory classes (songs, ballads, sagas, genres) independently of biological speciation, and the introduction of an ontologically-irreversible state-change operator into a multi-agent system. Empirical investigation of these layers is the subject of Phase 2 work, planned for the second half of 2026.

We invite collaboration, critique, and extension. The full mechanism registry, including the 30+ mechanisms not yet implemented, is included in the repository as an explicit invitation to other researchers who might want to take up specific entries.

Acknowledgements

This work was carried out by a single independent researcher with no institutional funding. Architectural design, the master mechanism registry, and methodological commitments are the author's. Code implementation was performed in iterative collaboration with several large language models of the Anthropic Claude family with additional assistance from Google Gemini and OpenAI ChatGPT. Empirical analyses are the author's.

References

- [1] T. S. Ray. An approach to the synthesis of life. In *Artificial Life II*, 1991.
- [2] C. Ofria and C. O. Wilke. Avida: A software platform for research in computational evolutionary biology. *Artificial Life*, 10(2):191–229, 2004.
- [3] L. Yaeger. Computational genetics, physiology, metabolism, neural systems, learning, vision, and behavior or PolyWorld: Life in a new context. In *Proceedings of the Artificial Life III Conference*, 1994.
- [4] A. Channon. Improving and still passing the ALife test: Component-normalised activity statistics classify evolution in Geb as unbounded. *Proceedings of Artificial Life VIII*, 2003.
- [5] H. Brighton and S. Kirby. Understanding linguistic evolution by visualizing the emergence of topographic mappings. *Artificial Life*, 12(2):229–242, 2006.
- [6] H. von Helmholtz. *Die Lehre von den Tonempfindungen als physiologische Grundlage für die Theorie der Musik*. Vieweg, 1863.

[TODO: Add 15–25 more references: Lehman & Stanley novelty search; recent Sakana AI publications on evolutionary model merging; Park et al. 2023 on Generative Agents; Lazaridou et al. on emergent communication; EGG framework; Whitehouse on modes of religiosity; Atran on cognitive science of religion; cross-cultural ritual studies; and any 2024–2026 ALife papers on dimensionality or niche emergence]